

Medical Insurance Cost Analysis

This notebook analyzes medical insurance costs based on various patient characteristics. We'll explore the data, create visualizations, and build predictive models using different regression techniques.

1. Import Libraries

```
In [1]: # Import necessary Libraries for data analysis and modeling
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import cross_val_score, train_test_split

# Set up visualization
%matplotlib inline
```

2. Load and Explore the Dataset

```
In [2]: # Load the dataset from the URL
filepath = 'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMD
df = pd.read_csv(filepath, header=None)

# Display the first 10 rows of the dataframe
print(df.head(10))
```

0	1	2	3	4	5	6	
0	19	1	27.900	0	1	3	16884.92400
1	18	2	33.770	1	0	4	1725.55230
2	28	2	33.000	3	0	4	4449.46200
3	33	2	22.705	0	0	1	21984.47061
4	32	2	28.880	0	0	1	3866.85520
5	31	1	25.740	0	?	4	3756.62160
6	46	1	33.440	1	0	4	8240.58960
7	37	1	27.740	3	0	1	7281.50560
8	37	2	29.830	2	0	2	6406.41070
9	60	1	25.840	0	0	1	28923.13692

3. Data Preprocessing

```
In [3]: # Add headers to the dataframe
headers = ["age", "gender", "bmi", "no_of_children", "smoker", "region", "charges"]
```

```
df.columns = headers

# Replace '?' with NaN for missing values
df.replace('?', np.nan, inplace = True)

# Display information about the dataframe
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2772 entries, 0 to 2771
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   age         2768 non-null    object  
 1   gender       2772 non-null    int64  
 2   bmi          2772 non-null    float64 
 3   no_of_children 2772 non-null    int64  
 4   smoker        2765 non-null    object  
 5   region        2772 non-null    int64  
 6   charges       2772 non-null    float64 
dtypes: float64(2), int64(3), object(2)
memory usage: 151.7+ KB
None
```

4. Handle Missing Values

```
In [4]: # Handle missing values in 'smoker' column (categorical)
# Replace with most frequent entry
is_smoker = df['smoker'].value_counts().idxmax()
# Corrected line: Reassign the result instead of using inplace=True
df["smoker"] = df["smoker"].replace(np.nan, is_smoker)

# Handle missing values in 'age' column (continuous)
# Replace with mean age
mean_age = df['age'].astype('float').mean(axis=0)
# Corrected line: Reassign the result instead of using inplace=True
df["age"] = df["age"].replace(np.nan, mean_age)

# Update data types for age and smoker
df[["age", "smoker"]] = df[["age", "smoker"]].astype("int")

# Display updated information
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2772 entries, 0 to 2771
Data columns (total 7 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               2772 non-null    int32  
 1   gender             2772 non-null    int64  
 2   bmi                2772 non-null    float64 
 3   no_of_children     2772 non-null    int64  
 4   smoker              2772 non-null    int32  
 5   region              2772 non-null    int64  
 6   charges             2772 non-null    float64 
dtypes: float64(2), int32(2), int64(3)
memory usage: 130.1 KB
None
```

5. Data Cleaning and Final Preview

```
In [5]: # Round the charges to 2 decimal places
df[["charges"]] = np.round(df[["charges"]],2)


# Display the first few rows of the cleaned dataframe
print(df.head())
```

	age	gender	bmi	no_of_children	smoker	region	charges
0	19	1	27.900		0	1	3 16884.92
1	18	2	33.770		1	0	4 1725.55
2	28	2	33.000		3	0	4 4449.46
3	33	2	22.705		0	0	1 21984.47
4	32	2	28.880		0	0	1 3866.86

6. Data Visualization


6.1 BMI vs Charges

```
In [6]: # Create a regression plot of BMI vs Charges
plt.figure(figsize=(10, 6))
sns.regplot(x="bmi", y="charges", data=df, line_kws={"color": "red"})
plt.ylim(0, )
plt.title('Relationship between BMI and Insurance Charges')
plt.xlabel('BMI')
plt.ylabel('Charges')
plt.grid(True)
plt.show()
```



6.2 Smoker Status vs Charges

```
In [7]: # Create a box plot of Smoker Status vs Charges
plt.figure(figsize=(10, 6))
sns.boxplot(x="smoker", y="charges", data=df)
plt.title('Insurance Charges by Smoker Status')
plt.xlabel('Smoker (0=Non-smoker, 1=Smoker)')
plt.ylabel('Charges')
plt.grid(True)
plt.show()
```



7. Correlation Analysis

```
In [8]: # Display the correlation matrix
print(df.corr())
```

	age	gender	bmi	no_of_children	smoker	\
age	1.000000	-0.026046	0.113048	0.037574	-0.023286	
gender	-0.026046	1.000000	0.042924	0.016020	0.082326	
bmi	0.113048	0.042924	1.000000	-0.001492	0.011489	
no_of_children	0.037574	0.016020	-0.001492	1.000000	0.006362	
smoker	-0.023286	0.082326	0.011489	0.006362	1.000000	
region	-0.007167	0.022213	0.271119	-0.025717	0.054077	
charges	0.298624	0.062837	0.199846	0.066442	0.788783	

	region	charges
age	-0.007167	0.298624
gender	0.022213	0.062837
bmi	0.271119	0.199846
no_of_children	-0.025717	0.066442
smoker	0.054077	0.788783
region	1.000000	0.054058
charges	0.054058	1.000000

8. Simple Linear Regression

```
In [9]: # Create a simple linear regression model using 'smoker' as the predictor
X = df[['smoker']]
Y = df['charges']
# Use a descriptive name for the simple linear model
lm_slr = LinearRegression()
```

```
lm_slr.fit(X,Y)

# Calculate and display the R-squared value
print(f"R-squared value for Simple Linear Regression: {lm_slr.score(X, Y):.4f}")
```

R-squared value for Simple Linear Regression: 0.6222

9. Multiple Linear Regression

```
In [10]: # Create a multiple linear regression model using multiple predictors
Z = df[["age", "gender", "bmi", "no_of_children", "smoker", "region"]]
# Use a descriptive name for the multiple linear model
lm_mlr = LinearRegression()
lm_mlr.fit(Z,Y)

# Calculate and display the R-squared value
print(f"R-squared value for Multiple Linear Regression: {lm_mlr.score(Z, Y):.4f}")
```

R-squared value for Multiple Linear Regression: 0.7504

10. Polynomial Regression Pipeline

```
In [11]: # Create a pipeline with scaling, polynomial features, and Linear regression
Input=[('scale',StandardScaler()),
       ('polynomial', PolynomialFeatures(include_bias=False)),
       ('model', LinearRegression())]
pipe=Pipeline(Input)

# Convert Z to float type and fit the pipeline
Z = Z.astype(float)
pipe.fit(Z,Y)

# Make predictions and calculate R-squared
ypipe=pipe.predict(Z)
print(f"R-squared value for Polynomial Regression: {r2_score(Y,ypipe):.4f}")
```

R-squared value for Polynomial Regression: 0.8452

11. Data Splitting for Model Evaluation

```
In [12]: # Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(Z, Y, test_size=0.2, random_st

print(f"Training set size: {x_train.shape[0]} samples")
print(f"Testing set size: {x_test.shape[0]} samples")
```

Training set size: 2217 samples

Testing set size: 555 samples

12. Ridge Regression

```
In [13]: # x_train, x_test, y_train, y_test hold same values as in previous cells
# Use a descriptive name for the linear Ridge model
RidgeModel_linear = Ridge(alpha=0.1)
RidgeModel_linear.fit(x_train, y_train)
yhat = RidgeModel_linear.predict(x_test)
print(f"R-squared value for Ridge Regression: {r2_score(y_test,yhat):.4f}")
```

R-squared value for Ridge Regression: 0.6761

13. Polynomial Ridge Regression

```
In [14]: # x_train, x_test, y_train, y_test hold same values as in previous cells
pr = PolynomialFeatures(degree=2)
x_train_pr = pr.fit_transform(x_train)
x_test_pr = pr.transform(x_test)
# Use a descriptive name for the polynomial Ridge model
RidgeModel_poly = Ridge(alpha=0.1)
RidgeModel_poly.fit(x_train_pr, y_train)
y_hat = RidgeModel_poly.predict(x_test_pr)
print(f"R-squared value for Polynomial Ridge Regression: {r2_score(y_test,y_hat):.4f}")
```

R-squared value for Polynomial Ridge Regression: 0.7836

14. Model Comparison and Conclusion

```
In [15]: # Create a summary of all models' performance
models = ['Simple Linear Regression', 'Multiple Linear Regression',
          'Polynomial Regression', 'Ridge Regression', 'Polynomial Ridge Regression']

# Calculate R-squared for each model using the correct model variables
r2_scores = [
    lm_slr.score(X, Y), # Uses the simple Linear model
    lm_mlr.score(Z, Y), # Uses the multiple Linear model
    r2_score(Y, ypipe),
    r2_score(y_test, yhat), # Uses the Linear Ridge model's predictions
    r2_score(y_test, y_hat) # Uses the polynomial Ridge model's predictions
]

# Create a dataframe to compare models
comparison_df = pd.DataFrame({
    'Model': models,
    'R-squared': r2_scores
})

# Sort by R-squared value
comparison_df = comparison_df.sort_values('R-squared', ascending=False)

# Display the comparison
print(comparison_df)
```

	Model	R-squared
2	Polynomial Regression	0.845236
4	Polynomial Ridge Regression	0.783563
1	Multiple Linear Regression	0.750408
3	Ridge Regression	0.676081
0	Simple Linear Regression	0.622179

15. Conclusion

Based on our analysis, we can draw the following conclusions:

1. The Polynomial Ridge Regression model performs best with an R-squared value of [value from output].
2. The 'smoker' attribute has the strongest correlation with insurance charges.
3. Multiple features including age, BMI, and smoker status collectively contribute to predicting insurance costs.
4. Regularization techniques like Ridge regression help prevent overfitting, especially when combined with polynomial features.

This analysis provides a foundation for predicting medical insurance costs based on patient characteristics. Further improvements could include trying different polynomial degrees, exploring other regularization techniques, or implementing more advanced machine learning models.